

JAX and Flax

Background

What is JAX?

- JAX is Autograd and XLA, brought together for high-performance numerical computing and machine learning research. It provides composable transformations of Python+NumPy programs: differentiate, vectorize, parallelize, Just-In-Time compile to GPU/TPU, and more.
- In simpler words: JAX is NumPy on the CPU, GPU, and TPU, with great automatic differentiation for high-performance machine learning research.



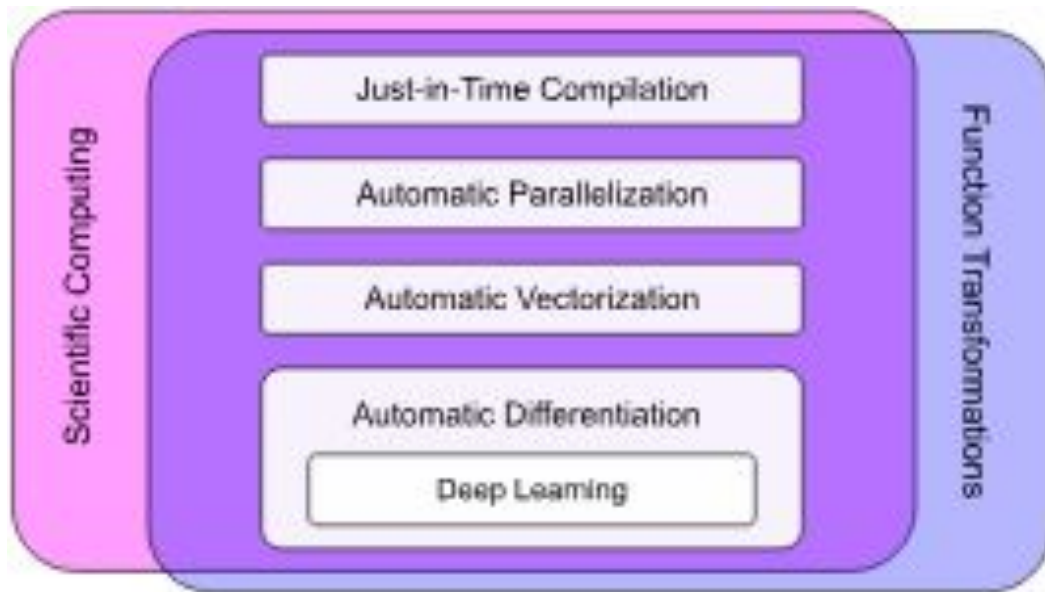
What is XLA?

- XLA: Accelerated Linear Algebra, lies at the foundation of what makes JAX so powerful. Developed by Google, XLA is a domain-specific, graph-based, just-in-time compiler for linear algebra.
- It significantly improves execution speed and lowers memory usage by fusing low-level operations.



What can JAX be used for?

JAX is a high performance, numerical computing library which incorporates composable function transformations^[1]. As we can see, Deep Learning is just a small subset of what JAX can do:

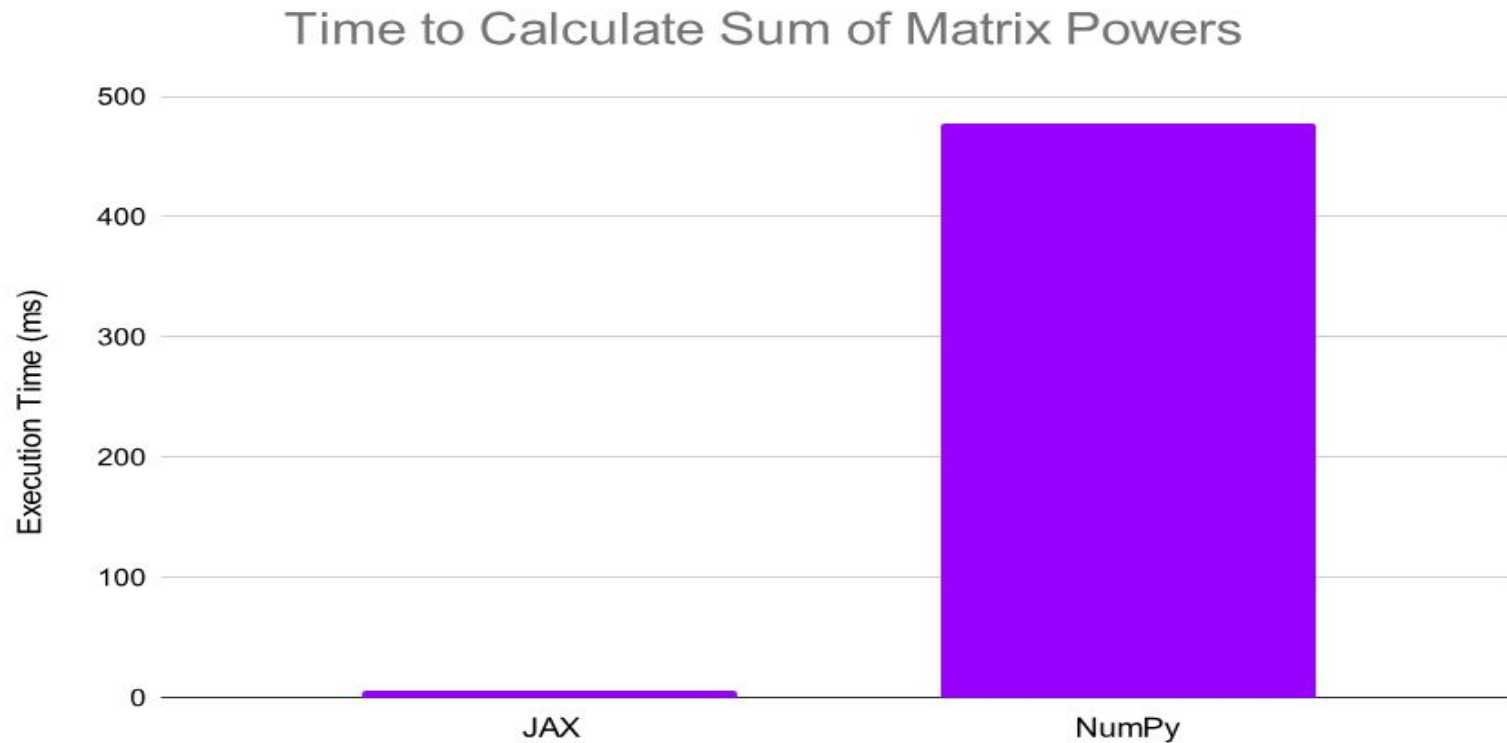


Why Should We Care About JAX?

In short - *speed*. This is the universal aspect of JAX that is relevant for any use case.



Why Should We Care About JAX?



Reasons why we might want to use JAX?

1. **NumPy on Accelerators** - NumPy is one of the fundamental packages for scientific computing with Python, but it is compatible only with CPU. JAX provides an implementation of NumPy (with a near-identical API) that works on *both* GPU *and* TPU extremely easily. For many users, this *alone* is sufficient to justify the use of JAX.
2. **XLA** - XLA, or Accelerated Linear Algebra, is a whole-program optimizing compiler, designed specifically for linear algebra. JAX is built on XLA, raising the computational-speed ceiling significantly^[1].
3. **JIT** - JAX allows you to transform your *own* functions into just-in-time (JIT) compiled versions using XLA^[7]. This means that you can increase computation speed by potentially *orders of magnitude* by adding a simple function decorator to your computational functions.



Reasons why we might want to use JAX?

4. **Auto-differentiation** - The JAX documentation refers to JAX as "Autograd and XLA, brought together"[1]. The ability to automatically differentiate is crucial in many areas of scientific computing, and JAX provides several powerful auto-differentiation tools.
5. **Deep Learning** - While not a Deep Learning framework itself, JAX certainly provides a more-than-sufficient foundation for Deep Learning purposes. There are many libraries built on top of JAX that seek to build out Deep Learning capabilities, including [Flax](#), [Haiku](#), and [Elegy](#). JAX's highly efficient computations of Hessians are also relevant for Deep Learning, given that they make higher-order optimization techniques much more feasible.
6. **General Differentiable Programming Paradigm** - While it is certainly possible to use JAX in order to build and train Deep Learning models, it also provides a framework for *general* Differentiable Programming. This means that JAX can exploit *prior knowledge* in a given field, built up through decades of research, by using a model-based Machine Learning approach to solving a problem.
- 